# TEK9010 - Exam prep summary

Sebastian G. Winther-Larsen

December 3, 2020

## 0: General concepts in multiagent systems

### Agents

Definition from Woolridge: "An agent is a computer system that is situated in some environment, and that is capable of autonomous action in this environment in order to meet its delegated objectives."

An agent is a computer system capable of independent action on behalf of its owner or user. Agents need skills and abilities to cooperate, coordinate and negotiate with each other on behalf of their users.

We deal with two different kinds of agents,

1. Reactive (or simple) agents can produce emergent properties usually modelled by Swarm Intelligence (SI). These can produce complex collective properties/performance.
2. Strategic (or intelligent) agents engage in social activities like cooperation, coordination, negotiation, etc usually described with Game Theory (GT). These engage in strategic interaction.

In some situations an agent-based solution would be appropriate.

1. The environment is open, or at least highly dynamic, uncertain or complex. In these settings autonomous agents might be the only solution.
2. Agents are natural metaphors for socities, organisations and businesses, as well as intelligent interfaces such as "expert assistants".
3. In istribution of data, control or expertise. When centralised solutions are difficult, like synchronisation of many autonomous databases.
4. When dealing with legacy software, one can wrap the legacy software in an agent layer.

### Multiagents systems

Multiagent systems consists of many agents interacting with each other through some network or sensor system. Definition from Woolridge: "Multiagent systems are systems of multiple interacting computing elements known as agents."

Multiagent systems is an appropriate software paradigm for modelling and building massive open and complex distributed systems. It is also a natural metaphor for artificial social systems. The research goal of multiagents systems is to connect behaviour on the microscopic scale with (often) emergent properties and effects on the macroscopic scale.

There are some key challenges associated with multiagent system desing.

1. The agent design problem on the microscopic level. How do we build agents that are capable of independent, autonomous action in order to successfully carry out tasks that we delegate to them?
2. The social desgin problem on the macroscopic level. How do we build agents that are capable of intercating with other agents in order to successfully carry out the tasks that we delegate to them, especially when the agents do no share common goals or intentions?

A truly succesful multiagent system makes an explicit connection between the autonomous micro level agents and the macro level modelling of the complex system.

Mutliagent system applications can be divided into two groups,

1. Distributed systems, where agents are processing nodes in a distributed multiagent system (emphasis on "multi").
2. Personal software assistants, where individual agents act as proactive assistants to users (emphasis on "individual").

## Swarm intelligence

Swarm Intelligence (SI) is the emergent collective intelligence of groups of simple agents, typically based on the social insect metaphor. This differs from Game Theory where agents are modellled as engaging in strategic interaction, typically inspired by utility-based models in the social sciences.

In terms of multiagents systems (MAS) there are some key similarities between the two paradigms. Both SI and GT model autonomous agents, which is the only requirement in MAS, and both SI and GT employ distributedness, i.e. there is no central control. These two points are actually the same. Some key differences are that SI use simple, reactive agents vs strategic, proactive, social agents in GT.

Emergence and stigmergy are two properties that are unique to SI.

## Stigmergy

Stigmergy is indirect, non-symbolic interaction mediated by the environment. Agents can use stigmergy to excange information by modifying the environment. This concet does not exist in game theoretical framework, because one assumes direct communication between agents.

**Emergence**

Emergence is an important in SI.

# 1: Non-cooperative game theory

## Different types of games

REMEMBER TO LOOK THROUGH A FEW EXAMPLES, including payoff matrices before exam.

### Cooperative / non-cooperative

A game is *cooperative* if the players are able to form binding commitments that are externally enforced. This external force can stem from contractual law or a some government. A game is *non-cooperative* if players cannot form or if all agreemetns need to be self-enforcing, e.g. through credible threats.

### Symmetric / asymmetric

A symmetric game is a game where the payoffs for playing a particular strategy depend only in the other strategies employed, not on who is playing them. If the identities of the players can be changed without changing the payoff to the strategies, the game is symmetric. Chicken, Prisoner's Dilemma and Stag Hunt are all symmetric games.

The most commonly studied asyummetric games are games where there are not identical strategy sets for both players. For instance, the Ultimatum Game and the Dictator Game have different strategies for each player. It is possible for a game to have identical stratgies for both (all) players, yet be asymmetric.

### Zero-sum games

Zero-sum games are a special case of constant-sum games in which choices by players can neither increase nor decrease the available resources. In zero-sum games, the total benefit to all players in the game, for every combination of strategies, always adds to zero. One can say that a player only benefits at equal expense of others.

Matching Pennies, poker and chess are all zero-sum games. Many games studied (like Prisoner's Dilemma), are not zero-sum games because the outcome has net results greater or less than zero.

Constant-sum games correspond to activities like theft and gambling, but not to the fundamental economic situation in which there are potential gains from trade. It is possible to transform any game in to (possibly asymmetric) zero-sum game by adding a dummy player whose losses compensate the players' net winnings.

### Simultaneous / sequential

*Simultaneous* games are games where both (all) players move simultaneously, or at least unaware of the other player(s)' earlier action. *Sequential* games or dynamic games, are games where players ahev some knowledge about earlier actions.

### Evolutionary game theory

Evolutionary game theory studies players who adjust their strategies over time according to rules that are not necessarily rational or farsighted. The fitness of individuals is not constant in evolutionary games, but depends on relative proportions or frequencies of the different phenotypes in the population.

The evolution of strategies over time is typically modelled as a Markov chain with a state variable such as the current strategy profile or how the game has been palyed in the recent past.

### Behavioural game theory

Behavioural game theory analyzes interactive strategic decisions and behaviour using methods of game theory, experimental economics and experimental psychology. Experiments include testing deviations from typical simplifications of economic theory such as the independence axiom and neglect altruism, fairness and framing effects.

## Self-interested agents

Agents have their own desires, preferences and beliefs. They seek to maximise their expected utility,

$$Ag_{\text{opt}} = \max \sum u(r) P(r|Ag, Env),$$

where $r$ is a "run". Whatever that means. This equation comes from the Tileworld example.

What is more, agent beliefs are modelled by information processes.

There is talk about a set of outcomes that an agent can experience,

$$\Omega = \{\omega_1, \omega_2, \dots\}.$$

## Utility

The utility of an agent is a valuable associated with an outcome,

$$u_i : \Omega \to \mathbb{R}$$

where $u_i$ is the utility of agent $i$, $\Omega$ is the set of possible outcomes, $\mathbb{R}$ is the set of real numbers.

Utility function is a way of representing an agent's preferences. Utility does not directly equate to money, but it helps to think of money. The utility value fo money is typically diminishing. An increase in utility from \$0 to \$1m is much greater than from \$500m to \$500m. You can make an inverse argument regarding debt.

Agents are able to rank outcomes by applying their utility function,

$$u_i(\omega_k) \geq u_i(\omega_l) \iff w_k \succeq_i w_l$$

means that agent $i$ prefers outcome $\omega_k$ to $\omega_l$, or is indefferent, whil

$$u_i(\omega_k) > u_i(\omega_l) \iff w_k \succ_i w_l,$$

means that agent $i$ strictly prefers outcome $\omega_k$ $\omega_l$.

Here are some important properties of preference ordering.

1. Reflexivity: $\omega \succeq_i \omega \ \forall \omega \in \Omega$
2. Transivity: $\omega_k \succeq_i \omega_l, \ \omega_l \succeq_i \omega_m \implies \omega_k \succeq_i \omega_m$
3. Comparability: Either $\omega_k \succeq_i \omega_l$ or $\omega_l \succeq_i \omega_k \ \forall \omega \in \Omega$.

## Strategic interaction

The environment is altered in simultaneous actions by agents. The basic idea of strategic interaction:

> What I do depends on what you do, and what you do depends on what I do, which we should both have taken into account in the first place.

One can assume that agents must act and that agents can not see other agents perform actions, in the scope of this course.

Mathematically we can write,

$$\tau : Ac_i \times Ac_j \to \Omega,$$

where $\tau$ is the state transformer function, $Ac_i$ is the action of agent $i$ and $\Omega$ is the set of outcomes.

The simplest strategic game conceivalbe consists of two agents, $i$ and $j$, with two actions available, Cooperate ($C$) or Defect ($D$). This gives us four different combinations of actions,

$$(C,C) \vee (C,D) \vee (D,C) \vee (D,D)$$

with four different outcomes,

$$\tau(C,C) = \omega_1, \ \tau(C,D) = \omega_2, \ \tau(D,C) = \omega_3, \ \tau(D,D) = \omega_4,$$

that is,

$$\Omega = \{\omega_1, \omega_2, \omega_4, \omega_4\}.$$

The agents $i$ and $j$ will value these outcomes according to their utility functions, $u_i$ and $u_j$.

Table 1: Payoff matrix. This is a game on strategic (or normal) form. The extensive form would be a game tree.

| i \ j | D | C |
|-------|-----------------|-----------------|
| D | $u_{i4}$, $u_{j4}$ | $u_{i3}$, $u_{j3}$ |
| C | $u_{i2}$, $u_{j2}$ | $u_{i1}$, $u_{j1}$ |

An interaction is said to be strictly competitive between agent $i$ and agent $j$ when

$$\omega \succ_i \omega' \wedge \omega' \succ_j \omega.$$

## Solution concepts for simple games

A set of concepts for solving games on strategic from;

1. Maximising social welfare
2. Pareto effiency
3. Dominant strategy
4. Nash equilibrium

Social welfare is given by the sum all all utilities for a particular outcome $\omega_i$,

$$sw(\omega_i) = \sum_{j \in Ag} u_j(\omega_i).$$

By choosing strategies that would maximise this function we have an outcome that gives the highest aggregated utility across agents.

A solution is Pareto efficient (or Pareto optimal) if no improvement is possible without making someone else worse off.

A strategy $s$ for agent $i$ is dominant if $s$ is the best response to all of the opposing agent $j$'s strateiges. There is no geuarnatee that such a solution exists.

**The Nash equilibrium**

If each agent has chosen a strategy, based on what they other player(s) is doing, and no agent can increase its own expected payoff by changing its strategy while the other players keep theirs unchanged, then the current set of strategy choices constitutes a Nash equilibrium.

The two strategies $s_i$ and $s_j$ of agents $i$ and $j$ are in Nash equilibrium,

1. if agent $i$ plays $s_i$, player $j$ can do no better than playing $s_j$ and
2. if agent $j$ plays $s_j$, player $i$ can do no better than playing $s_i$.

$s_i$ ans $s_j$ are best repsonses to the other and no agent regrets its strategy choice.

A Nash equilibrium is *weak* if a change of strategy would yield the exact same output. If there is only one Nash "equilibrating" strategy, we have a *strict*

7

Nash equilibrium. A game can have a pure-strategy or a mixed-strategy Nash equilibrium. A *pure strategy* provides a complete definition of how a player will play a game. A *mixed strategy* is an assignment of a probability to each pure strategy. This allows for a player to randomly select a pure strategy. A *totally mixed* strategy is a mixed strategy in which the agent assigns a strictly positive probability to every pure strategy.

Nash's theorem states that every game in which every agent has a finite set of possibilities has a Nash equilibrium in mixed strategies.

## Prisoner's Dilemma

Two man are collectively charged with a crime and held in separate cells. They have no way of communicating with each other or making any kind of agreement. The two men are told that if one of the confesses the crime and the other does not, the confessor will be freed, and the other will be jailed for 3 years. If both confess to the crime, then each will be jailed for 2 years. If neither prisoner confesses they will be jailed for 1 year.

Table 2: Payoff-matrix for agent $i$ in PD.

|   | C | D |
|---|---|---|
| C | -1 | -3 |
| D | 0 | -2 |

This is a dilemma, because the Nash equilibrium is also the worst social outcome.

You should also add some utilities etc!!

Some important real-world examples of PD are the "Tragedy of the commons", like grazing livestock, overfishing of the seas and capacity bandwidth on the Internet. Nuclear weapons treaties and international environmental agreements as well. Perhaps any game with positive externalities?

## Repeated PD

Repeating PD over several rounds will increase the change of cooperation, because of the threat of "punishment" by defecting in susequent rounds and because loss of utility can be amortized over several rounds.

If there are infinite rounds in the repeated PD game, cooperation is rational due to threat of defection. If the number of rounds is a fixed number, it is rational to defect in the last round. That in turn makes it rational to defect in the second-to-last round. Go figure. If there is a probability for playing another round it is rational to cooperate if the probability of one more round is large enough compared to the payoffs.

# 9: Swarm Robotics 1

## What is swarm robotics?

Quick answer: swarm intelligence applied to robotics.

There is no explicit definition of a *swarm* in literature. A swarm is defined via its behavior.

The *size of a swarm* is defined by what it is not: "not as large as to be dealt with statistical averages" and "not as small as to be dealt with as a few-body problem". The size of a swarm $N$ is

$$10^2 < N << 10^2 3,$$

not Avagadro-large.

Swarm robotics is "the study of how a large number of relatively simple physically embodied agents can be designed such that a desired collective behavior emerges from local interactions among agents and between agents and the environment", according to Dorigo and Sahin. But! A swarm is not necessarily

There are some key features. The fact that local interactions between agents and the environment should be possible requires robots to have local sensing and probably also communication capabilities. In fact, (local) communication is considered a key feature of swarms.

Collaboration is required to go beyond a mere paralllisation in swarm a swarm system. We want to go beyond the performance of simple parallelisation. Think of some clearning task with each robot cleaning a small assigned area.

## Swarm performance.

Some keywords are *contention* or *inference* and (lack of) *coherency*, given by parameters $\alpha$ and $\beta$, repectively. The robots need to share limited resources and communicate. This is difficult.

In the contexst of swarm robotics we can interpret contention as interference between robots due to shared resouces, such as an entrance to a base station or generally space. Collision avoidance is a waiting loop because the shared resource *space* is currently not available. This can be compared to an airplane flying in a holding pattern because the resource "runway" is currently in use and should certainly not be shared. Incoherency, in turn, can be interpreted as inconsistencies or overhead due to limited communication of imformation or due to imperfect synchrony.

The univeral scalability law is important,

$$R(N) = \frac{N}{1 + \alpha(N - 1) + \beta N(N - 1)}.$$

Its inverntor, Gunther, identifies four qualitatively different situations,

1. If contention and lack of coherency are negligible, then we get "equal bang for the buck" and have a linear speedup ($\alpha = 0$, $\beta = 0$),
2. If there is a cost for sharing resources in the form of contention, then we have a sublinear speedup ($\alpha > 0$, $\beta = 0$),
3. If ther is an increased negative influence due to contention, then the speedup clearly levels off ($\alpha >> 0$, $\beta = 0$).
4. If in additon there is also an increased influence of incoherence, then there exists a peak speed up and for bigger system sizes the speedup decreases ($\alpha >> 0$, $\beta > 0$).



Figure 1: Gunther's Universal law of Computational Scalability

One can identify some "regions" of performance; super-linear, sub-linear, optimal and inference. As more agents are added, performance starts to increase (sub-/super-)linearly, then we get to an optimum after a while. After that comes the inference region.

In parallel computing, superlinear speedups can occur due to some interplay between problem size per computing unit and available memory. For example, if the problem can be divied into pieces that fit completely into a CPUs cache, then one can observe a considerable speedup. In swarm robotics, superlinear performance increases occur due to qualitatively different collaboration modes that are accessible with increasing swarm size as seen in bucket brigades.

It is possible for a system-wide deadlock to occur in a swarm robotics system. For instance with a very high swarm denisty, such that all robots permanently try to avoid collisons resulting in zero performance.

## Modelling swarms (as a series of mappings).

A swarm system of size $N$ in 2D space can be described by the state vecotr,

$$\gamma = (r_1, r_2, \ldots, r_N, v_1, v_2, \ldots, v_N, s_1, s_2, \ldots, s_N),$$

wher $r_i$ are the positions, $v_i$ are the velocities and $s_i$ are the discrete states for agents $i \in [1, N]$. We denote the configuration space by $\Gamma$, i.e. $\gamma \in \Gamma$, with $\dim \Gamma = 2N + 2N + N = 5N$.

This is a large space to keep track off, and we can only observe one sequence, $\gamma_t, \gamma_{t+1}, \gamma_{t+2}, \ldots$ at a time for a specific initial state $\gamma_0$. What we ideally need instead is to understand how the system operates in gemeral for any setup. We need to omit certain parts, s.t. we obtain a different definition of a configuration, $\phi \in \Phi$ with $\dim \Phi << \dim \Gamma$. We seek a mapping,

$$f : \Gamma \mapsto \Phi.$$

In a discrete time model, we also need two update rules, $g : \Gamma \mapsto \Gamma$ and $h : \Phi \mapsto \Phi$. We now have the following functionality,

$$f(\gamma_t) = \phi_t, \ g(\gamma_t) = \gamma_{t+1}, \ h(\phi_t) = \phi_{t+1}.$$

The following requirement should hold (we want this),

$$h(f(\gamma_t)) = f(g(\gamma_t)),$$

that is, the modelling abstractions implemented by $f$ should be chosen carefulle sucht that the model update rule $h$ is able to predict the correct model configuration for the next time step.

We would also like an inverse map $f^{-1} : \Phi \mapsto \Gamma$ that reverses the model abstractions and rebuilds the configuration $\gamma$ of the real system from the model configuration $\phi$. However, typically $f^{-1}$ cannot usefully be defined, because $f$ is surjective, that is, we can have $\gamma_1, \gamma_2 \in \Gamma$ with $\gamma_1 \neq \gamma_2$ and $f(\gamma_1) = f(\gamma_2)$ due to the reduction in dimensionality caused by $f$.

Why do we need models in Swarm Robotics? Quote from Schweitzer:

> To gain insight into the interplay between microscopic interactions and macroscopicfeatures, it is important to find a level of description that, on the one hand, considers specificfeatures of the system and is suitable for reflecting the origination of new qualities, but, onthe other hand, is not flooded with microscopic details. (. . . ) A commonly accepted theory of agent systems that also allows analytical investigations is, however, still pending because of the diversity of the various models invented for particular applications.

An extreme example of an extreme model abstraction,

$$f(\gamma) = \frac{|\{s_i | s_i = A\}|}{N} = \phi,$$

where $\dim \Phi = 1$. How does the update rule $h$ look like? Non-trivial!

In swarm robotics we are still on the search for an appropriate general modelling technique.

## When are rate equations appropriate?

Rate equations are appropriate when there is no special information of interest in the full state of the system. We are only interested in macroeconomics properties This is typically when we are working with concentrations, for example. Such systems are typically inspired by chemical systems.

## The Langevin equation.

$$\dot{\mathbf{R}}(t) = \mathbf{A}(\mathbf{R}(t), t) + B(\mathbf{R}(t), t)\mathbf{F}(t),$$

where $\mathbf{R}$ is the position of an agent, $\mathbf{F}$ is a stochastic process (e.g. white noise), $\mathbf{A}$ describes and scales the agent's behaviour and $B$ describes and scales non-deterministic behaviour. A possible choice for $\mathbf{A}$ is a gradient descent in a potential field,

$$\mathbf{A}(\mathbf{R}(t), t) = \nabla P(\mathbf{R}(t), t).$$

## The Focker-Planck equation.

The ~ is a PDE describing the temporal dynamics of a probability density. This density describes in the original physical conetext the probability of finding the particle within a certain area. It is the macroscopically corresponding piece to the microscopic approach described by the Langevin equation. It was originannly used in physics for modelling Brownian motion with drift, describing diffusion processes in thermodynamics.

So, it looks like this,

$$\frac{\partial \rho(\mathbf{r}, t)}{\partial t} = \nabla(\mathbf{A}(\mathbf{r}, t)\rho(\mathbf{r}, t)) + \frac{1}{2}Q\nabla^2(B^2(\mathbf{r}, t)\rho(\mathbf{r}, r)),$$

where $\rho$ is a (probability) density for a single particle at position $\mathbf{r}$ and time $t$. We interpret it as the robot density of all coexisiting robots of the swarm. By integrating over an area $W$,

$$s(t) = \int_{\mathbf{r} \in W} \rho(\mathbf{r}, t),$$

we get the expected fration for the swarm within that area at time $t$. The first term in the Fokker-Planck equation is a non-stochastic drift term and the second term is a diffusion term.