

TEK9010 - Exam prep summary

Sebastian G. Winther-Larsen

December 4, 2020

0: General concepts in multiagent systems

Agents

Definition from Woolridge: “An agent is a computer system that is situated in some environment, and that is capable of autonomous action in this environment in order to meet its delegated objectives.”

An agent is a computer system capable of independent action on behalf of its owner or user. Agents need skills and abilities to cooperate, coordinate and negotiate with each other on behalf of their users.

We deal with two different kinds of agents,

1. Reactive (or simple) agents can produce emergent properties usually modelled by Swarm Intelligence (SI). These can produce complex collective properties/performance.
2. Strategic (or intelligent) agents engage in social activities like cooperation, coordination, negotiation, etc usually described with Game Theory (GT). These engage in strategic interaction.

In some situations an agent-based solution would be appropriate.

1. The environment is open, or at least highly dynamic, uncertain or complex. In these settings autonomous agents might be the only solution.
2. Agents are natural metaphors for societies, organisations and businesses, as well as intelligent interfaces such as “expert assistants”.
3. In distribution of data, control or expertise. When centralised solutions are difficult, like synchronisation of many autonomous databases.
4. When dealing with legacy software, one can wrap the legacy software in an agent layer.

Multiagents systems

Multiagent systems consists of many agents interacting with each other through some network or sensor system. Definition from Woolridge: “Multiagent systems are systems of multiple interacting computing elements known as agents.”

Multiagent systems is an appropriate software paradigm for modelling and building massive open and complex distributed systems. It is also a natural metaphor for artificial social systems. The research goal of multiagents systems is to connect behaviour on the microscopic scale with (often) emergent properties and effects on the macroscopic scale.

There are some key challenges associated with multiagent system desing.

1. The agent design problem on the microscopic level. How do we build agents that are capable of independent, autonomous action in order to successfully carry out tasks that we delegate to them?
2. The social desgin problem on the macroscopic level. How do we build agents that are capable of intercatng with other agents in order to successfully carry out the tasks that we delegate to them, especially when the agents do no share common goals or intentions?

A truly succesful multiagent system makes an explicit connection between the autonomous micro level agents and the macro level modelling of the complex system.

Mutliagent system applications can be divided into two groups,

1. Distributed systems, where agents are processing nodes in a distributed multiagent system (emphasis on “multi”).
2. Personal software assistants, where individual agents act as proactive assistants to users (emphasis on “individual”).

Swarm intelligence

Swarm Intelligence (SI) is the emergent collective intelligence of groups of simple agents, typically based on the social insect metaphor. This differs from Game Theory where agents are modelled as engaging in strategic interaction, typically inspired by utility-based models in the social sciences.

In terms of multiagents systems (MAS) there are some key similarities between the two paradigms. Both SI and GT model autonomous agents, which is the only requirement in MAS, and both SI and GT employ distributedness, i.e. there is no central control. These two points are actually the same. Some key differences are that SI use simple, reactive agents vs strategic, proactive, social agents in GT.

Emergence and stigmergy are two properties that are unique to SI.

Stigmergy

Stigmergy is indirect, non-symbolic interaction mediated by the environment. Agents can use stigmergy to exchang information by modifying the environment. This concet does not exist in game theoretical framework, because one assumes direct communication between agents.

Emergence

Emergence is an important in SI.

1: Non-cooperative game theory

Different types of games

REMEMBER TO LOOK THROUGH A FEW EXAMPLES, including payoff matrices before exam.

Cooperative / non-cooperative

A game is *cooperative* if the players are able to form binding commitments that are externally enforced. This external force can stem from contractual law or a some government. A game is *non-cooperative* if players cannot form or if all agreements need to be self-enforcing, e.g. through credible threats.

Symmetric / asymmetric

A symmetric game is a game where the payoffs for playing a particular strategy depend only in the other strategies employed, not on who is playing them. If the identities of the players can be changed without changing the payoff to the strategies, the game is symmetric. Chicken, Prisoner's Dilemma and Stag Hunt are all symmetric games.

The most commonly studied asymmetric games are games where there are not identical strategy sets for both players. For instance, the Ultimatum Game and the Dictator Game have different strategies for each player. It is possible for a game to have identical strategies for both (all) players, yet be asymmetric.

Zero-sum games

Zero-sum games are a special case of constant-sum games in which choices by players can neither increase nor decrease the available resources. In zero-sum games, the total benefit to all players in the game, for every combination of strategies, always adds to zero. One can say that a player only benefits at equal expense of others.

Matching Pennies, poker and chess are all zero-sum games. Many games studied (like Prisoner's Dilemma), are not zero-sum games because the outcome has net results greater or less than zero.

Constant-sum games correspond to activities like theft and gambling, but not to the fundamental economic situation in which there are potential gains from trade. It is possible to transform any game in to (possibly asymmetric) zero-sum game by adding a dummy player whose losses compensate the players' net winnings.

Simultaneous / sequential

Simultaneous games are games where both (all) players move simultaneously, or at least unaware of the other player(s)' earlier action. *Sequential* games or dynamic games, are games where players have some knowledge about earlier actions.

Evolutionary game theory

Evolutionary game theory studies players who adjust their strategies over time according to rules that are not necessarily rational or farsighted. The fitness of individuals is not constant in evolutionary games, but depends on relative proportions or frequencies of the different phenotypes in the population.

The evolution of strategies over time is typically modelled as a Markov chain with a state variable such as the current strategy profile or how the game has been played in the recent past.

Behavioural game theory

Behavioural game theory analyzes interactive strategic decisions and behaviour using methods of game theory, experimental economics and experimental psychology. Experiments include testing deviations from typical simplifications of economic theory such as the independence axiom and neglect altruism, fairness and framing effects.

Self-interested agents

Agents have their own desires, preferences and beliefs. They seek to maximise their expected utility,

$$Ag_{\text{opt}} = \max \sum u(r)P(r|Ag, Env),$$

where r is a “run”. Whatever that means. This equation comes from the Tileworld example.

What is more, agent beliefs are modelled by information processes.

There is talk about a set of outcomes that an agent can experience,

$$\Omega = \{\omega_1, \omega_2, \dots\}.$$

Utility

The utility of an agent is a valuable associated with an outcome,

$$u_i : \Omega \rightarrow \mathbb{R}$$

where u_i is the utility of agent i , Ω is the set of possible outcomes, \mathbb{R} is the set of real numbers.

Utility function is a way of representing an agent's preferences. Utility does not directly equate to money, but it helps to think of money. The utility value for money is typically diminishing. An increase in utility from \$0 to \$1m is much greater than from \$500m to \$500m. You can make an inverse argument regarding debt.

Agents are able to rank outcomes by applying their utility function,

$$u_i(\omega_k) \geq u_i(\omega_l) \iff w_k \succeq_i w_l$$

means that agent i prefers outcome ω_k to ω_l , or is indifferent, while

$$u_i(\omega_k) > u_i(\omega_l) \iff w_k \succ_i w_l,$$

means that agent i strictly prefers outcome ω_k to ω_l .

Here are some important properties of preference ordering.

1. Reflexivity: $\omega \succeq_i \omega \forall \omega \in \Omega$
2. Transitivity: $\omega_k \succeq_i \omega_l, \omega_l \succeq_i \omega_m \implies \omega_k \succeq_i \omega_m$
3. Comparability: Either $\omega_k \succeq_i \omega_l$ or $\omega_l \succeq_i \omega_k \forall \omega \in \Omega$.

Strategic interaction

The environment is altered in simultaneous actions by agents. The basic idea of strategic interaction:

What I do depends on what you do, and what you do depends on what I do, which we should both have taken into account in the first place.

One can assume that agents must act and that agents can not see other agents perform actions, in the scope of this course.

Mathematically we can write,

$$\tau : Ac_i \times Ac_j \rightarrow \Omega,$$

where τ is the state transformer function, Ac_i is the action of agent i and Ω is the set of outcomes.

The simplest strategic game conceivable consists of two agents, i and j , with two actions available, Cooperate (C) or Defect (D). This gives us four different combinations of actions,

$$(C, C) \vee (C, D) \vee (D, C) \vee (D, D)$$

with four different outcomes,

$$\tau(C, C) = \omega_1, \tau(C, D) = \omega_2, \tau(D, C) = \omega_3, \tau(D, D) = \omega_4,$$

that is,

$$\Omega = \{\omega_1, \omega_2, \omega_4, \omega_4\}.$$

The agents i and j will value these outcomes according to their utility functions, u_i and u_j .

An interaction is said to be strictly competitive between agent i and agent j when

$$\omega \succ_i \omega' \wedge \omega' \succ_j \omega.$$

Table 1: Payoff matrix. This is a game on strategic (or normal) form. The extensive form would be a game tree.

$i \setminus j$	D	C
D	u_{i4}, u_{j4}	u_{i3}, u_{j3}
C	u_{i2}, u_{j2}	u_{i1}, u_{j1}

Solution concepts for simple games

A set of concepts for solving games on strategic form;

1. Maximising social welfare
2. Pareto efficiency
3. Dominant strategy
4. Nash equilibrium

Social welfare is given by the sum all all utilities for a particular outcome ω_i ,

$$sw(\omega_i) = \sum_{j \in Ag} u_j(\omega_i).$$

By choosing strategies that would maximise this function we have an outcome that gives the highest aggregated utility across agents.

A solution is Pareto efficient (or Pareto optimal) if no improvement is possible without making someone else worse off.

A strategy s for agent i is dominant if s is the best response to all of the opposing agent j 's strategies. There is no guarantee that such a solution exists.

The Nash equilibrium

If each agent has chosen a strategy, based on what the other player(s) is doing, and no agent can increase its own expected payoff by changing its strategy while the other players keep theirs unchanged, then the current set of strategy choices constitutes a Nash equilibrium.

The two strategies s_i and s_j of agents i and j are in Nash equilibrium,

1. if agent i plays s_i , player j can do no better than playing s_j and
2. if agent j plays s_j , player i can do no better than playing s_i .

s_i and s_j are best responses to the other and no agent regrets its strategy choice.

A Nash equilibrium is *weak* if a change of strategy would yield the exact same output. If there is only one Nash “equilibrating” strategy, we have a *strict* Nash equilibrium. A game can have a pure-strategy or a mixed-strategy Nash equilibrium. A *pure strategy* provides a complete definition of how a player will play a game. A *mixed strategy* is an assignment of a probability to each pure strategy. This allows for a player to randomly select a pure strategy. A *totally mixed* strategy is a mixed strategy in which the agent assigns a strictly positive probability to every pure strategy.

Nash’s theorem states that every game in which every agent has a finite set of possibilities has a Nash equilibrium in mixed strategies.

Prisoner’s Dilemma

Two men are collectively charged with a crime and held in separate cells. They have no way of communicating with each other or making any kind of agreement. The two men are told that if one of the confesses the crime and the other does not, the confessor will be freed, and the other will be jailed for 3 years. If both confess to the crime, then each will be jailed for 2 years. If neither prisoner confesses they will be jailed for 1 year.

Table 2: Payoff-matrix for agent i in PD.

	C	D
C	-1 -3	
D	0 -2	

This is a dilemma, because the Nash equilibrium is also the worst social outcome.

You should also add some utilities etc!!

Some important real-world examples of PD are the “Tragedy of the commons”, like grazing livestock, overfishing of the seas and capacity bandwidth on the Internet. Nuclear weapons treaties and international environmental agreements as well. Perhaps any game with positive externalities?

Repeated PD

Repeating PD over several rounds will increase the chance of cooperation, because of the threat of “punishment” by defecting in subsequent rounds and because loss of utility can be amortized over several rounds.

If there are infinite rounds in the repeated PD game, cooperation is rational

due to threat of defection. If the number of rounds is a fixed number, it is rational to defect in the last round. That in turn makes it rational to defect in the second-to-last round. Go figure. If there is a probability for playing another round it is rational to cooperate if the probability of one more round is large enough compared to the payoffs.

2: Voting

Social choice theory

In a basic setting we have a set of agents,

$$Ag = \{1, 2, \dots, N\},$$

where the number of agents $|Ag| = N$ is assumed to be finite and odd, in order to break ties. We also have a set of possible outcomes,

$$\Omega = \{\omega_1, \omega_2, \dots\},$$

where $|\Omega| = k$ is the number of possible outcomes. In a pairwise election $k = 2$ and in a general election $k \geq 2$.

For each (voting) agent i we have a preference ordering,

$$\bar{\omega}_i = (\omega_3, \omega_1, \dots, \omega_k).$$

Preference aggregation is the fundamental problem in social choice theory. How do we combine the different agents' preference orderings in order to derive a group decision? More specific, how do we generate a social preference order over possible outcomes?

We need a *social welfare function* that takes the voter preferences and produces a social preferences order,

$$f : \Pi(\Omega) \times \dots \Pi(\Omega) \times \rightarrow \Pi(\Omega).$$

We will also use $\omega \succ^* \omega'$ to indicate that ω is ranked over ω' in the social outcome.

Sometimes, we will be concerned with slightly simpler settings, in which we are not concerned with obtaining an intire ordering, but just one of the possible candidates. We use the term *social choice function* for such a mapping,

$$f : \Pi(\Omega) \times \dots \Pi(\Omega) \times \rightarrow \Omega.$$

Plurality voting

In plurality voting, every voter submits their preference order and the winner is the otucome ranked first most times. This is the simplest and best known voting procedure. Plurality is straightforward to implement and easy to understand by the voters. However, plurality is vulnerable to strategic manipulation and tactical votion. Additionally it reveals Condorcet's paradox.

Simple majority voting

If we only have two outcomes to choose between, then plurality is just *simple majority voting*. In this case, we simply ask voters to select one of the two outcomes, and the one that gets the majority of votes is the winner. Simple majority voting is not very easy to manipulate. In reality there are often more than two possible outcomes..

Sequential majority election

In a *sequential majority election* a series of plurality elections are conducted to determine a winner. The idea is that a pair of outcomes will face each other in pairwise election, and the winner will then go on to the next election. These sequential pairwise elections can be arranged into linear order of ballots, or a binary tree election.

Suppose we are voting over four outcomes $\omega_1, \omega_2, \omega_3$ and ω_4 . We might choose the following order or *agenda* for the election,

$$\omega_2, \omega_3, \omega_4, \omega_1.$$

We could have the first election between ω_2 and ω_3 , then the winner would face ω_4 and so on. This is a *linear order*. If we structured the election as *balanced binary tree*, we would have simultaneous elections, e.g. the winner between ω_2 and ω_3 would face the winner between ω_4 and ω_1 .

Bear in mind that the ultimate outcome is generally sensitive to the election agenda, or voting order.

Borda count

For each of the possible candidate outcomes $|\Omega| = k$, we have a numerical value counting the strength of opinion in favour of this candidate. If an outcome ω_j appears first in the preference order, we increment the count for ω_j by $k - 1$, we then increment the count for the next outcome by $k - 1$ and so on until the final outcome in the preference order is incremented by 0. We proceed with this process until all preference orders have been considered, and then simply order the outcomes Ω by their counts, in descending order.

The Borda count (BC) for outcome ω_j is given by

$$BC_{\omega_j} = \sum_{i=1}^N k - \text{rank}(\bar{\omega}_i(\omega_j)),$$

where $k = |\Omega|$ is the number of possible outcomes, N is the number of voters and $\bar{\omega}_i$ is the social preference order of voter i .

Slater rule

READ THROUGH THE example in the book closely (p.260).

The Slater rank is used for breaking cycles in majority graphs. The Slater rank for a social ordering is how many edges must be “flipped” in the cyclic majority graph to produce that particular social order.

The Slater rule is to choose the social ordering that minimises the disagreement between the majority graph and the social choice, i.e. the order with lowest Slater rank number.

Dictatorship

A social welfare function is said to be a dictatorship if

$$f(\bar{\omega}_1, \dots, \bar{\omega}_i, \dots, \bar{\omega}_N) = \bar{\omega}_i,$$

where $\bar{\omega}_i$ is the preference order of voter i . Consequently, in dictatorship, the social outcome is only dependent on voter i .

Arrow’s theorem

Assuming voters have three or more distinct alternatives, there exist *no* ranked voting electoral system that can convert the ranked preferences of individuals into a social preference order while at the same time also meet a set of specific “desirable” criteria,

1. Unrestricted domain
2. Pareto efficiency
3. Independence of irrelevant alternatives (IIA)
4. Non-dictatorship

The *unrestricted domain* condition states that all preferences of all voters are allowed, meaning that the preference order of the voters should be complete and the social preference order should be deterministic.

The *Pareto condition* states that there is no other outcome that makes one voter better off without making any other voter worse off, i.e. if all voters $\omega \succ \omega' \implies \omega \succ^* \omega'$. This condition is satisfied for Plurality, Borda and dictatorship, but not for sequential majority election.

IIA: The social preference between outcome $\omega \succ^* \omega'$ depends only on the individual preferences between $\omega \succ \omega'$. This means that the ranking of all the other outcomes, not changing the relative ranking of individual ranking of $\omega \succ \omega'$ should affect the social ranking of $\omega \succ^* \omega'$. Dictatorship satisfies this criterion, but Plurality, Borda and sequential majority election do not.

If the non-dictatorship condition is dropped as a criterion, then a dictatorship satisfies Arrow’s theorem!

Majority graph

A *majority graph* is a directed graph constructed from voter preferences. An outcome is a *possible winner* if there is some agenda which would result in that outcome being the overall winner. An outcome is called the *Condorcet winner* if it is the overall winner for every possible agenda.




Figure 1: Majority graph where every outcome is a possible winner

Properties of a majority graph

1. Completeness: For any two outcomes ω_i and ω_j , we must have either ω_i defeat ω_j or ω_j defeat ω_i .
2. Asymmetry: If ω_i defeats ω_j then ω_j cannot defeat ω_i .
3. Irreflexivity: ω_i will never defeat itself.

Condorcet's paradox

There are scenarios where no matter what outcome we choose, a majority of voters will be unhappy with the outcome. This will depend on preference ordering of voters.

Suppose we have three outcomes, $\Omega = \{\omega_1, \omega_2, \omega_3\}$, and three voters, $Ag = \{1, 2, 3\}$, with the following preference orderings,

$$\bar{\omega}_1 = (\omega_1, \omega_2, \omega_3) \iff \omega_1 \succ \omega_2 \succ \omega_3,$$

$$\bar{\omega}_2 = (\omega_3, \omega_1, \omega_2) \iff \omega_3 \succ \omega_1 \succ \omega_2,$$




Figure 2: Majority graph where the outcome represented by the node in the middle is a Condorcet winner

$$\bar{\omega}_3 = (\omega_2, \omega_3, \omega_1) \iff \omega_1 \succ \omega_3 \succ \omega_1,$$

Selecting ω_1 would mean two thirds of voters would rather prefer ω_3 , and the same goes for any of the other outcomes.

Tactical voting and strategic manipulation

Tactical voting is a way of strategically misrepresenting ones preferences in order to bring about a more preferred outcome. Suppose you are in a voting district that leans more towards the Conservative candidate, with Labour running third in terms of percentage support. You personally favour Labour, with a Liberal Democrat second and a Conservative third. With a simple plurality voting, placing the Labour candidate first in your ranking any well be wasted. If you instead rank the Liberal candidate first, you might be able to get the Liberal candidate elected, and bring about an outcome that you prefer over the election of a Conservative candidate.

A social welfare function is manipulable if there exists some $\hat{\omega}'_i$ such that

$$f(\hat{\omega}_1, \dots, \hat{\omega}'_i, \dots, \hat{\omega}_N) \succ_i f(\hat{\omega}_1, \dots, \hat{\omega}_i, \dots, \hat{\omega}_N),$$

where $\hat{\omega}_i$ is the preference order of voter i . This means that the social outcome could be improved for some voter i by unilaterally misrepresenting i 's preference order.

The Gibbard-Satterthwaite theorem

Assuming voters have three or more distinct alternatives, according to *the Gibbard-Satterthwaite theorem*, in general there exist *no* voting protocol, except for dictatorship, that is non-manipulable. However, strategic manipulation might be hard to compute and uncertainties could make manipulation strategies more difficult to obtain.

Second-order Copeland Scheme

For a voting procedure to be “easy to compute”, we mean that the function f can be implemented by an algorithm that runs in time polynomial in the number of voters and candidates. For a voting procedure to be “easy to manipulate”, we mean that if it is for a voter i to obtain a more preferred outcome by declaring a preference order ω'_i rather than its true preference ω_i , then such a ω'_i can be computed in polynomial time. Now here’s a question: are there non-dictatorial voting procedures that are easy to compute and that satisfy the Pareto condition, but are *not* easy to manipulate? Yes! The *Second-order Copeland* satisfies these requirements. It also satisfies the Condorcet winner property.

3: Cooperative Game Theory

The game we have studied most is PD. The following feature of PD prevents cooperation,

- Binding agreements are not possible
- Utility is given directly to individuals as a result of individual action

The first feature means that agents cannot trust one another’s promises. They just make their decisions based solely on the information they have about strategies and individual pay-offs, will try to maximise their own utility, and will assume that all other agents will be trying to do the same thing. With binding agreement, PD would be no D at all.

The second feature says that an individual need not be concerned with collective utility; it need only to try to maximise its own utility and assume that everybody else will do likewise.

Cooperative game theory attempts to answer two basic questions,

1. Which coalition will be formed by self-interested rational agents?
2. How is the utility divided among members in this coalition?

Formal description of a cooperative game

We have a set of agents

$$Ag = \{1, 2, \dots, N\},$$

that can form coalitions C, C', C_1, \dots that are subsets of Ag , e.g. $C_1 \subset Ag$. We if $C = Ag$, C is the *Grand Coalition*, and if $C = \{i\}$, C is called a *singleton* coalition.

A cooperative game (or a coalitional game) is a pair,

$$G = \langle Ag, v \rangle,$$

where Ag is the set of agents and $v : 2^{Ag}$ is the characteristic function. The characteristic function assigns a numerical value to all possible coalitions. $v(C) = k$ means that the value k is assigned to coalition C and distributed among its members.

The game itself is completely silent about how this utility should be distributed among coalition members. Coalition members have to agree amongst themselves on how to divide the “pay cheque”. The origin of the characteristic function for any given scenario is generally regarded as not being a concern in cooperative game theory. It is assumed to simply be “given”.

We can identify three key issues that have to be addressed with respect to cooperative games. These issues are referred to as the *cooperation lifecycle*.

1. Coalition structure generation: which coalitions will form, or rather, which coalitions are *stable*.
2. Solve the optimisation problem for each coalition. Solving the joint problem of the coalition, i.e. finding the best way to maximise the utility of the coalition itself.
3. Dividing the value of the solution for each coalition. The *Shapley value* is a “fair” distribution of utility amongst coalition members.

The core

The core of a coalitional game is the set of feasible distributions of pay-off to members of a coalition that no subcoalition can reasonably object to. Formally, a coalition $C' \subseteq Ag$ objects to an outcome $\mathbf{x} \langle x_1, \dots, x_N \rangle$ of the coalition C if there exist some outcome $\mathbf{x}' \langle x'_1, \dots, x'_N \rangle$ in C' such that $x'_i > x_i \forall i \in C'$. The outcomes satisfy the characteristic function v ,

$$v(C) = \sum_{i \in C} x_i.$$

A coalition C is stable if the core is non-empty.

Suddenly outcomes are written with the letter \mathbf{x} .

Example. Consider a game $G = \langle Ag, v \rangle$ where $Ag = \{1, 2\}$, $v(\{1\}) = 5$, $v(\{2\}) = 5$ and $v(\{1, 2\}) = 20$. The grand coalition of G is stable for several outcomes,

$$v(\{1, 2\}) = \langle 5, 15 \rangle = 20,$$

$$v(\{1, 2\}) = \langle 6, 14 \rangle = 20,$$

$$v(\{1, 2\}) = \langle 7, 13 \rangle = 20,$$

and so on. If for instance,

$$v(\{1, 2\}) = \langle 19, 1 \rangle = 20 \rightarrow v(\{2\}) = 5,$$

agent 2 objects.

The question remains - how do we divide utility amongst coalition members? A good idea is to *divide surplus according to contribution*.

The Shapley value

The Shapley value is based on the idea that agents should get the average marginal contribution it makes to a coalition, estimated over all possible positions that it would enter the coalition. The Shapley value is the unique value that satisfies the “fairness” axioms.

If $\mu_i(C)$ is the amount that i adds by joining $C \subseteq Ag \setminus \{i\}$ the *marginal contribution* of i to C ,

$$\mu_i(C) = v(C \cup \{i\}) - v(C).$$

If $\mu_i(C) = v(\{i\})$, then there is no synergy or added value from i joining C , since the amount i adds is exactly what i would earn on its own.

The three axioms are known as,

1. Symmetry: Agents that make the same contribution should get the same payoff. i and j are said to be interchangeable if $\mu_i(C) = \mu_j(C)$ for $C \subseteq Ag \setminus \{i, j\}$ and $sh_i = sh_j$.
2. Dummy player: Player who don't have any synergy with any coalition should get only what they can earn on their own. If $\mu_i = v(\{i\})$ for $C \subseteq Ag \setminus \{i\}$ then $Sh_i = v(\{i\})$.
3. Additivity: If you combine several games, the value that a player gets should be the sum of the value it gets in the individual games. A player does not gain or lose by playing more than once (There is a formal representation of this on p. 275 in MAS).

The Shapley value of agent i is given by,

$$Sh_i = \frac{1}{|Ag|!} \sum_{o \in \prod(Ag)} \mu_i(C_i(o)),$$

where $\prod(Ag)$ is the set of all possible orderings of coalition C , o is a specific ordering of a coalition and μ_i is the marginal contribution of agent i .

Example. We have a cooperative game $G = \langle Ag, v \rangle$ where $Ag = \{1, 2, 3\}$, $v(\{1\}) = v(\{2\}) = v(\{3\}) = 5$, $v(\{1, 2\}) = v(\{1, 3\}) = 10$, $v(\{2, 3\}) = 20$ and $v(\{1, 2, 3\}) = 25$.

Permutations of $\{1, 2, 3\}$	Marginal contributions of $\{1, 2, 3\}$
(1, 2, 3)	$\rightarrow (5, 5, 15)$
(1, 3, 2)	$\rightarrow (5, 15, 5)$
(2, 1, 3)	$\rightarrow (5, 5, 15)$
(2, 3, 1)	$\rightarrow (5, 5, 15)$
(3, 1, 2)	$\rightarrow (5, 15, 5)$
(3, 2, 1)	$\rightarrow (5, 15, 5)$

From the table, we can compute the Shapley value to $(30, 60, 60)/6$.

The Shapley value is hard to compute and represent for a large number of players (exponential in N). Is it possible to represent different coalition permutations in a more succinct and tractable way?

- Induced subgraph representation is succinct but not complete
- Marginal contribution nets are complete and succinct

Marginal contribution net

A marginal contribution net is an extension to the induced subgraph. The characteristic function of a game is represented by a set of rules,

$$rs_C = \{\phi \rightarrow x \in rs \mid C \models \phi\},$$

where $\phi \rightarrow x$ is a rule, C is a coalition and we write $C \models \phi$ to indicate that the rule ϕ applies to C . The characteristic function v_{rs} associated with the rule set rs is

$$v_{rs}(C) = \sum_{\phi \rightarrow x \in rs_C} x,$$

where $\phi \rightarrow x$ is a rule in the rule set rs_C .

Look at this in book and/or slides!

Simple games

A simple coalitional game is one in which every coalition gets a value of either 1 (“winning”) or 0 (“losing”). Simple games are important because they model yes/no voting systems in which a proposal is pitted against the status quo. Formally we have a pair,

$$Y = \langle Ag, W \rangle,$$

where $Ag = \{1, 2, \dots, N\}$ is a set of N agents or voters and $W \subseteq 2^{Ag}$ is the set of winning coalitions such that $v(W) = 1$ and $v(C \notin W) = 0$.

Some properties of simple games:

1. Non-triviality. There are some winning coalitions, but not all coalitions are winning. Formally $\emptyset \subset W \subset 2^{Ag}$.

2. Monotonicity. If C wins, then every superset of C also wins. Formally if $C_1 \subseteq C_2$ and $C_1 \in W$ then $C_2 \in W$.
3. Zero-sum. If a coalition C wins, then agents outside C do not win. Formally if $C \in W$ then $Ag \setminus C \notin W$.
4. Empty coalition loses. $\emptyset \notin W$
5. Grand coalition wins. $Ag \in W$.

Explicitly listing all coalitions will be exponential in number of players.

Weighted voting games are a natural extension of simple games possibly reducing number of players in these games.

Weighted voting games

Weighted voting games is a more concise way of representing many simple games. For instance, instead of representing 100 US senators explicitly, we represent the different “blocks” of voters and evaluate those against the criteria of winning, called the quota. This greatly reduces the number of players/voters.

For each agent $i \in Ag$ we define a weight w_i , and we define an overall *quota* q . A coalition C is then winning if the sum of their weights exceeds the quota,

$$v(C) = \begin{cases} 1 & \text{if } \sum_{i \in C} w_i > q \\ 0 & \end{cases},$$

where $c(C)$ is the characteristic function of coalition C . In a simple majority voting process $q = \frac{|Ag|+1}{2}$.

The weighted voting game can be written on the form

$$\langle q; w_1, w_2, \dots, w_N \rangle.$$

The Shapley value calculates the power of voters in the game. How? It is NP-hard.

The Core is much easier to compute. How?

A special version of exists called the k -weighted voting game.

Coalitional structure formation

Central forming of coalition

1. All nodes owned by a single designer
2. Maximising social welfare.

The problem is to find the coalition structure that maximises the aggregated outcome. This is a search through a partition of the overall set of agents Ag into mutually disjoint coalitions.

Example using $Ag = \{1, 2, 3\}$. The possible coalitions are

$$\{1\}, \{2\}, \{3\}, \{1, 2\}, \{1, 3\}, \{2, 3\}, \{1, 2, 3\},$$

and possible coalition structures are,

$$\{1, 2, 3\}, \{1\}\{2, 3\}, \{2\}\{1, 3\}, \{3\}\{1, 2\}, \{1\}\{2\}\{3\}.$$

The social choice is

$$CS^* = \arg$$

4: Auctions

5: Bargaining

6: Arguing

7: Classical Swarm Intelligence

8: Response Threshold

9: Swarm Robotics 1

What is swarm robotics?

Quick answer: swarm intelligence applied to robotics.

There is no explicit definition of a *swarm* in literature. A swarm is defined via its behavior.

The *size of a swarm* is defined by what it is not: “not as large as to be dealt with statistical averages” and “not as small as to be dealt with as a few-body problem”. The size of a swarm N is

$$10^2 < N << 10^23,$$

not Avagadro-large.

Swarm robotics is “the study of how a large number of relatively simple physically embodied agents can be designed such that a desired collective behavior emerges from local interactions among agents and between agents and the environment”, according to Dorigo and Sahin. But! A swarm is not necessarily

There are some key features. The fact that local interactions between agents and the environment should be possible requires robots to have local sensing and probably also communication capabilities. In fact, (local) communication is considered a key feature of swarms.

Collaboration is required to go beyond a mere parallelisation in a swarm system. We want to go beyond the performance of simple parallelisation. Think of some cleaning task with each robot cleaning a small assigned area.

Swarm performance.

Some keywords are *contention* or *inference* and (lack of) *coherency*, given by parameters α and β , respectively. The robots need to share limited resources and communicate. This is difficult.

In the context of swarm robotics we can interpret contention as interference between robots due to shared resources, such as an entrance to a base station or generally space. Collision avoidance is a waiting loop because the shared resource *space* is currently not available. This can be compared to an airplane flying in a holding pattern because the resource “runway” is currently in use and should certainly not be shared. Incoherency, in turn, can be interpreted as inconsistencies or overhead due to limited communication of information or due to imperfect synchrony.

The universal scalability law is important,

$$R(N) = \frac{N}{1 + \alpha(N - 1) + \beta N(N - 1)}.$$

Its inventor, Gunther, identifies four qualitatively different situations,

1. If contention and lack of coherency are negligible, then we get “equal bang for the buck” and have a linear speedup ($\alpha = 0, \beta = 0$),
2. If there is a cost for sharing resources in the form of contention, then we have a sublinear speedup ($\alpha > 0, \beta = 0$),
3. If there is an increased negative influence due to contention, then the speedup clearly levels off ($\alpha \gg 0, \beta = 0$).
4. If in addition there is also an increased influence of incoherence, then there exists a peak speed up and for bigger system sizes the speedup decreases ($\alpha \gg 0, \beta > 0$).

One can identify some “regions” of performance; super-linear, sub-linear, optimal and inference. As more agents are added, performance starts to increase (sub-/super-)linearly, then we get to an optimum after a while. After that comes the inference region.

In parallel computing, superlinear speedups can occur due to some interplay between problem size per computing unit and available memory. For example, if the problem can be divided into pieces that fit completely into a CPU’s cache, then one can observe a considerable speedup. In swarm robotics, superlinear performance increases occur due to qualitatively different collaboration modes that are accessible with increasing swarm size as seen in bucket brigades.

It is possible for a system-wide deadlock to occur in a swarm robotics system.

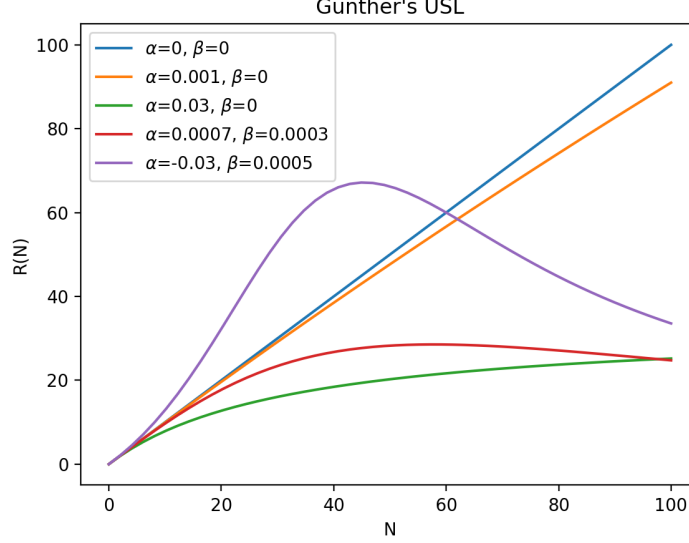


Figure 3: Gunther's Universal law of Computational Scalability

For instance with a very high swarm density, such that all robots permanently try to avoid collisions resulting in zero performance.

Modelling swarms (as a series of mappings).

A swarm system of size N in 2D space can be described by the state vector,

$$\gamma = (r_1, r_2, \dots, r_N, v_1, v_2, \dots, v_N, s_1, s_2, \dots, s_N),$$

where r_i are the positions, v_i are the velocities and s_i are the discrete states for agents $i \in [1, N]$. We denote the configuration space by Γ , i.e. $\gamma \in \Gamma$, with $\dim \Gamma = 2N + 2N + N = 5N$.

This is a large space to keep track of, and we can only observe one sequence, $\gamma_t, \gamma_{t+1}, \gamma_{t+2}, \dots$ at a time for a specific initial state γ_0 . What we ideally need instead is to understand how the system operates in general for any setup. We need to omit certain parts, s.t. we obtain a different definition of a configuration, $\phi \in \Phi$ with $\dim \Phi \ll \dim \Gamma$. We seek a mapping,

$$f : \Gamma \mapsto \Phi.$$

In a discrete time model, we also need two update rules, $g : \Gamma \mapsto \Gamma$ and $h : \Phi \mapsto \Phi$. We now have the following functionality,

$$f(\gamma_t) = \phi_t, \quad g(\gamma_t) = \gamma_{t+1}, \quad h(\phi_t) = \phi_{t+1}.$$

The following requirement should hold (we want this),

$$h(f(\gamma_t)) = f(g(\gamma_t)),$$

that is, the modelling abstractions implemented by f should be chosen carefully such that the model update rule h is able to predict the correct model configuration for the next time step.

We would also like an inverse map $f^{-1} : \Phi \mapsto \Gamma$ that reverses the model abstractions and rebuilds the configuration γ of the real system from the model configuration ϕ . However, typically f^{-1} cannot usefully be defined, because f is surjective, that is, we can have $\gamma_1, \gamma_2 \in \Gamma$ with $\gamma_1 \neq \gamma_2$ and $f(\gamma_1) = f(\gamma_2)$ due to the reduction in dimensionality caused by f .

Why do we need models in Swarm Robotics? Quote from Schweitzer:

To gain insight into the interplay between microscopic interactions and macroscopic features, it is important to find a level of description that, on the one hand, considers specific features of the system and is suitable for reflecting the origination of new qualities, but, on the other hand, is not flooded with microscopic details. (...) A commonly accepted theory of agent systems that also allows analytical investigations is, however, still pending because of the diversity of the various models invented for particular applications.

An extreme example of an extreme model abstraction,

$$f(\gamma) = \frac{|\{s_i | s_i = A\}|}{N} = \phi,$$

where $\dim \Phi = 1$. How does the update rule h look like? Non-trivial!

In swarm robotics we are still on the search for an appropriate general modelling technique.

When are rate equations appropriate?

Rate equations are appropriate when there is no special information of interest in the full state of the system. We are only interested in macroeconomic properties. This is typically when we are working with concentrations, for example. Such systems are typically inspired by chemical systems.

The Langevin equation.

$$\dot{\mathbf{R}}(t) = \mathbf{A}(\mathbf{R}(t), t) + B(\mathbf{R}(t), t)\mathbf{F}(t),$$

where \mathbf{R} is the position of an agent, \mathbf{F} is a stochastic process (e.g. white noise), \mathbf{A} describes and scales the agent's behaviour and B describes and scales non-deterministic behaviour. A possible choice for \mathbf{A} is a gradient descent in a potential field,

$$\mathbf{A}(\mathbf{R}(t), t) = -\nabla P(\mathbf{R}(t), t).$$

The Fokker-Planck equation.

The \sim is a PDE describing the temporal dynamics of a probability density. This density describes in the original physical context the probability of finding the particle within a certain area. It is the macroscopically corresponding piece to the microscopic approach described by the Langevin equation. It was originally used in physics for modelling Brownian motion with drift, describing diffusion processes in thermodynamics.

So, it looks like this,

$$\frac{\partial \rho(\mathbf{r}, t)}{\partial t} = \nabla(\mathbf{A}(\mathbf{r}, t)\rho(\mathbf{r}, t)) + \frac{1}{2}Q\nabla^2(B^2(\mathbf{r}, t)\rho(\mathbf{r}, t)),$$

where ρ is a (probability) density for a single particle at position \mathbf{r} and time t . We interpret it as the robot density of all coexisting robots of the swarm. By integrating over an area W ,

$$s(t) = \int_{\mathbf{r} \in W} \rho(\mathbf{r}, t),$$

we get the expected fraction for the swarm within that area at time t . The first term in the Fokker-Planck equation is a non-stochastic drift term and the second term is a diffusion term.

10: Swarm Robotics 2